Analysis Report

horizontal_diffusion_gpu(int, int, int, int, int, int, int, int, int, int, int, int, int, int, int, int, int, int, int, int, int, int, int, int, int, char, float*, float*, float*, float*, float*, float*, float*, float*, float*, float*, float, float*, float, float)

Duration	10.158 ms (10,158,193 ns)
Grid Size	[14,39,1]
Block Size	[32,8,1]
Registers/Thread	63
Shared Memory/Block	6.641 KiB
Shared Memory Requested	16 KiB
Shared Memory Executed	16 KiB
Shared Memory Bank Size	4 B

| [0] GeForce GT 730	
Compute Capability	2.1
Max. Threads per Block	1024
Max. Shared Memory per Block	48 KiB
Max. Registers per Block	32768
Max. Grid Dimensions	[65535, 65535, 65535]
Max. Block Dimensions	[1024, 1024, 64]
Max. Warps per Multiprocessor	48
Max. Blocks per Multiprocessor	8
Number of Multiprocessors	2
Multiprocessor Clock Rate	1.4 GHz
Concurrent Kernel	true
Max IPC	4
Threads per Warp	32
Global Memory Bandwidth	22.4 GB/s
Global Memory Size	2 GiB
Constant Memory Size	64 KiB
L2 Cache Size	128 KiB
Memcpy Engines	1
PCIe Generation	2
PCIe Link Rate	5 Gbit/s
PCIe Link Width	4

1. Compute, Bandwidth, or Latency Bound

The first step in analyzing an individual kernel is to determine if the performance of the kernel is bounded by computation, memory bandwidth, or instruction/memory latency. The results below indicate that the performance of kernel "horizontal_diffusion_gpu" is most likely limited by compute. You should first examine the information in the "Compute Resources" section to determine how it is limiting performance.

1.1. Kernel Performance Is Bound By Compute

For device "GeForce GT 730" the kernel's memory utilization is significantly lower than its compute utilization. These utilization levels indicate that the performance of the kernel is most likely being limited by computation on the SMs.

2. Compute Resources

GPU compute resources limit the performance of a kernel when those resources are insufficient or poorly utilized. Compute resources are used most efficiently when instructions do not overuse a function unit. The results below indicate that compute performance may be limited by overuse of a function unit.

## 2.1. GPU Utilization Is Limited By Function Unit Usage

Different types of instructions are executed on different function units within each SM. Performance can be limited if a function unit is over-used by the instructions executed by the kernel. The following results show that the kernel's performance is potentially limited by overuse of the following function units: Arithmetic.
Load/Store - Load and store instructions for local, shared, global, constant, etc. memory.
Arithmetic - All arithmetic instructions including integer and floating-point add and multiply, logical and binary operations, etc.
Control-Flow - Direct and indirect branches, jumps, and calls.
Texture - Texture operations.
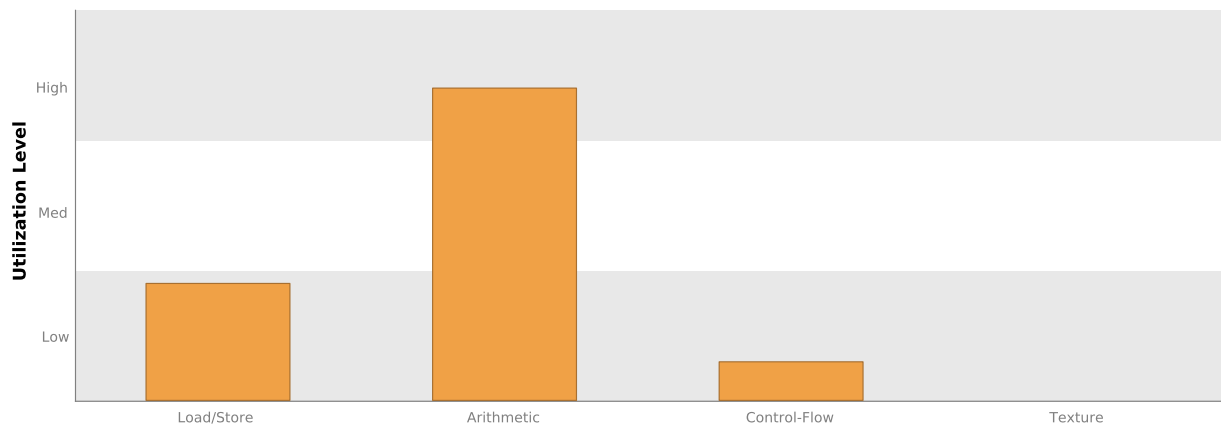
2.2. Instruction Execution Counts

The following chart shows the mix of instructions executed by the kernel. The instructions are grouped into classes and for each class the chart shows the percentage of thread execution cycles that were devoted to executing instructions in that class. The "Inactive" result shows the thread executions that did not execute any instruction because the thread was predicated or inactive due to divergence.
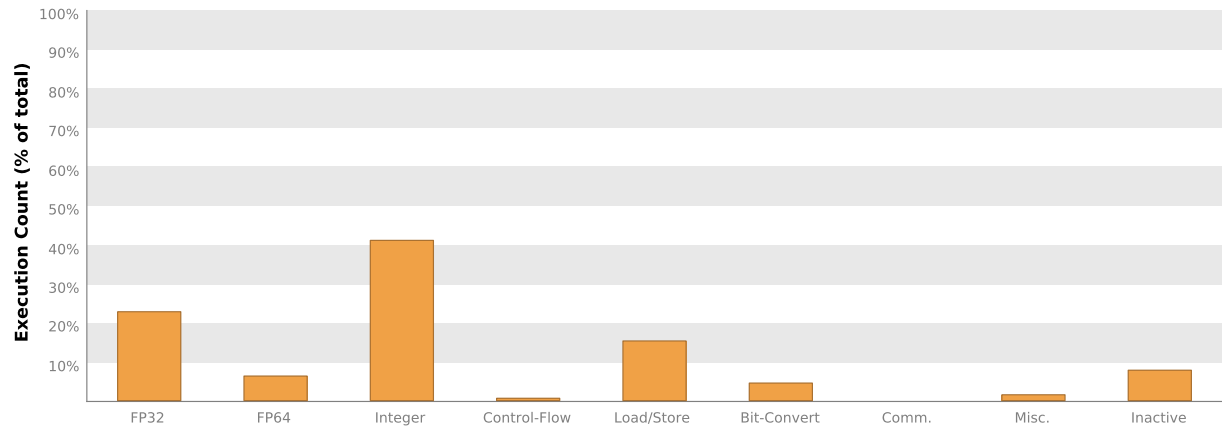
2.3. Floating-Point Operation Counts

The following chart shows the mix of floating-point operations executed by the kernel. The operations are grouped into classes and for each class the chart shows the percentage of thread execution cycles that were devoted to executing operations in that class. The results do not sum to 100% because non-floating-point operations executed by the kernel are not shown in this chart.

3. Memory Bandwidth

Memory bandwidth limits the performance of a kernel when one or more memories in the GPU cannot provide data at the rate requested by the kernel. The results below indicate that the kernel is limited by the bandwidth available to the device memory.

3.1. Memory Bandwidth And Utilization

The following table shows the memory bandwidth used by this kernel for the various types of memory on the device. The table also shows the utilization of each memory type relative to the maximum throughput supported by the memory.

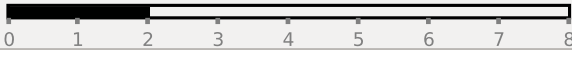|	Transactions	Bandwidth	Utilization
L1/Shared Memory			
Local Loads	0	0 B/s	
Local Stores	0	0 B/s	
Shared Loads	66336	827.994 MB/s	
Shared Stores	82920	1.035 GB/s	
Global Loads	3392660	42.347 GB/s	
Global Stores	270072	2.052 GB/s	
Atomic	0	0 B/s	
L1/Shared Total	3811988	46.262 GB/s	Idle Low Medium High Max
L2 Cache			
L1 Reads	4083604	12.743 GB/s	
L1 Writes	657636	2.052 GB/s	
Texture Reads	0	0 B/s	
Atomic	0	0 B/s	
Total	4741240	14.795 GB/s	Idle Low Medium High Max
Texture Cache			
Reads	0	0 B/s	Idle Low Medium High Max
Device Memory			
Reads	3214475	10.031 GB/s	
Writes	630945	1.969 GB/s	
Total	3845420	11.999 GB/s	Idle Low Medium High Max
System Memory			
[PCIe configuration: Gen2 x4, 5 Gbit/s]			
Reads	9	28.084 kB/s	Idle Low Medium High Max
Writes	0	0 B/s	Idle Low Medium High Max
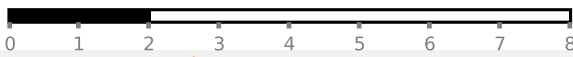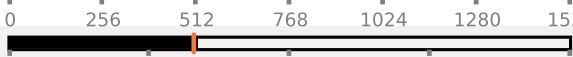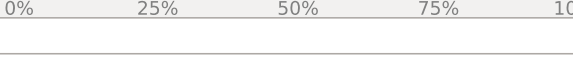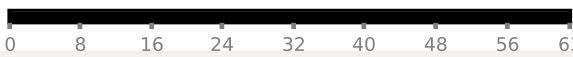
4. Instruction and Memory Latency

Instruction and memory latency limit the performance of a kernel when the GPU does not have enough work to keep busy. The performance of latency-limited kernels can often be improved by increasing occupancy. Occupancy is a measure of how many warps the kernel has active on the GPU, relative to the maximum number of warps supported by the GPU. Theoretical occupancy provides an upper bound while achieved occupancy indicates the kernel's actual occupancy. The results below indicate that occupancy can be improved by reducing the number of registers used by the kernel.

4.1. GPU Utilization Is Limited By Register Usage

The kernel uses 63 registers for each thread (16128 registers for each block). This register usage is likely preventing the kernel from fully utilizing the GPU. Device "GeForce GT 730" provides up to 32768 registers for each block. Because the kernel uses 16128 registers for each block each SM is limited to simultaneously executing 2 blocks (16 warps). Chart "Varying Register Count" below shows how changing register usage will change the number of blocks that can execute on each SM.

Optimization: Use the -maxrregcount flag or the __launch_bounds__ qualifier to decrease the number of registers used by each thread. This will increase the number of blocks that can execute on each SM. On devices with Compute Capability 5.2 turning global cache off can increase the occupancy limited by register usage.

Variable	Achieved	Theoretical	Device Limit	Grid Size: [14,39,1] (546 blocks) Block Size: [32,8,1] (256 thread
Occupancy Per SM				
Active Blocks		2	8	
Active Warps	15.55	16	48	
Active Threads		512	1536	
Occupancy	32.4%	33.3%	100%	
Warps				
Threads/Block		256	1024	
Warps/Block		8	32	
Block Limit		6	8	
Registers				
Registers/Thread		63	63	
Registers/Block		16384	32768	
Block Limit		2	8	
Shared Memory				
Shared Memory/Block		6800	16384	
Block Limit		2	8	

4.2. Occupancy Charts

The following charts show how varying different components of the kernel will impact theoretical occupancy.

Varying Block Size

Warps Per SM

16 @ 256

Threads Per Block

Varying Register Count

16 @ 63

Registers Per Thread

Varying Shared Memory Usage

16 @ 6k

Shared Memory Per Block (bytes)